

CIFAR-100

Kevin Lopez
California State University Long Beach
Long Beach, Ca
Kevin.LopezChavez01@student.csulb.edu

I. INTRODUCTION

This report details the development and evaluation of three Convolutional Neural Network (CNN) models trained purely on the CIFAR-100 dataset. It covers the experimental setup (in Section II), multiple model evaluations (in Section III), the model selection (in Section IV), the full retraining process (in Section V), and a comparison with top models in [papers with code](#).

II. EXPERIMENTAL SETUP

A. Data Splitting

The CIFAR-100 dataset is made up of 50,000 training images and 10,000 test images. The training dataset was further divided into:

- **Validation Set:** $\frac{1}{5}$ of entire training set = 10,000 images
- **Sub Train:** $Entire_training_set - \frac{1}{5} \times Entire_training_set = 40,000$ images

This split was done using PyTorch's `random_split` function.

B. Data Transformation

To make the models generalize better and reduce overfitting, I used data transformation and applied them to the entire training data. Transformations are the following:

- Random Horizontal Flip
- Random Crop with padding
- Color Jittering (brightness, contrast, saturation, hue)
- Random Rotation

These transformations allowed me to achieve a 10 percent improvement compared to doing it without. The same test was performed without data transformations and got 50 % percent accuracy for all three models. While I have not attached the .ipynb file, it is available upon request.

III. MODEL EVALUATION WITH RAY TUNE

To start, I developed **8 different CNN architectures**. Each architecture was trained using Ray Tune for hyper-parameter optimization and model selection. Ray Tune was configured to run **16 models in parallel** using asynchronous hyperband scheduling, keeping the top 50% of the models based on validation performance during training. In total, **60 models** were evaluated using ray tune and were run up to a maximum of 30 epochs.

The **Spatial Output Size Formula** formula was used to calculate the width and height of the output of convolution and pooling layers.

$$\text{Output Size} = \left\lfloor \frac{\text{Input Size} + 2 \times \text{Padding} - \text{Kernel Size}}{\text{Stride}} \right\rfloor + 1$$

All the models were designed to take in a 32x32x3 image with an output of 100 to handle the 100 classes in CIFAR100. The loss function used was multi-class cross-entropy to handle loss in the 100 different classes. $L = -\sum_{i=1}^N y_i \log(\hat{y}_i)$ It was necessary to add Batch Normalization and dropout to prevent the models from overfitting. I often got overfitting models whenever I trained the models without normalization and dropout.

1) *Summary of components used within the 60 models:*
Below is a summary of the different components used to evaluate the models

• Activation Functions

- ReLU: $f(x) = \max(0, x)$
- LeakyReLU: $f(x) = \begin{cases} x & \text{if } x \geq 0 \\ \alpha x & \text{if } x < 0 \end{cases}$ where α is small

• Pooling Layers

- I often used `kernel_size=2`, `stride=2`, `padding=0`, which allowed me to reduce the size of the CONV layers by half.
- MaxPool2d: Takes maximum value in each window
- AvgPool2d: Takes average value in each window

• Convolutional Layers

- Kernel sizes: 1x1, 2x2, 3x3, 4x4, 5x5, 7x7
- Stride=1, padding=1. (it often kept the same dimensions)

• To avoid overfit

- BatchNorm2d: Normalizes activations across batch dimension
- Dropout: Rates vary from 0.1 to 0.3

• Fully Connected Layers

- Used to extract the outputs from the Convolution layers and get the 100 different classes.
- At the output layer, the Cross EntropyLoss was used (Multiclass Loss Function).

• Optimizers:

- **Adam:** $lr \in [10^{-4}, 10^{-1}]$, $B1 \in [0.8, 0.99999]$, $B2 \in [0.8, 0.99999999]$
- **RMSProp:** $lr \sim \log U(10^{-4}, 10^{-1})$, $\alpha \in [0.8, 1]$, momentum $\in [0, 0.5]$

- **AdaGrad:** $lr \sim \log U(10^{-3}, 10^{-1})$,
initial_accumulator_value $\in [0.0, 0.3]$,
 $lr_decay \in [0, 0.1]$
- **SGD:** $lr \sim \log U(10^{-5}, 0.1)$, momentum $\in [0.0, 0.4]$, dampening $\in [0, 0.4]$

Figure 1 shows the best three performances (training and validation loss) while evaluating with ray tune.

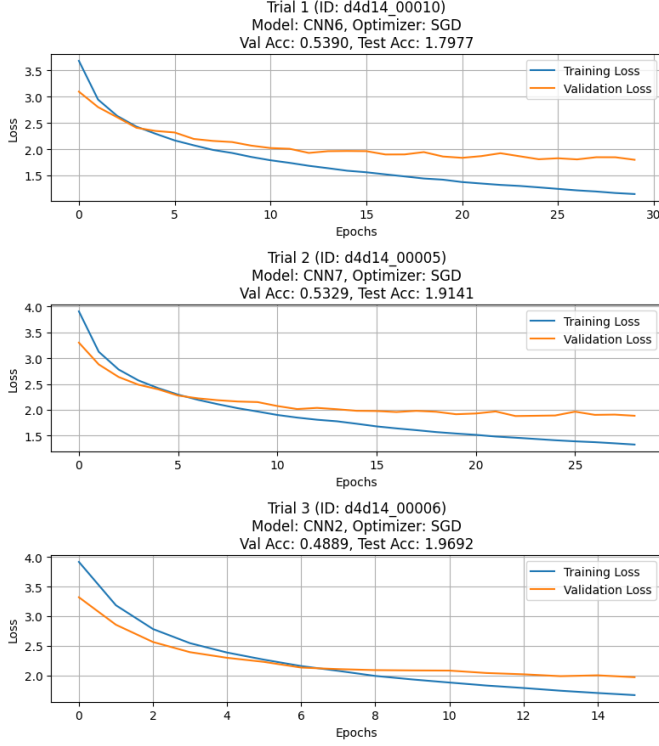


Fig. 1: Model Selection Training/Validation Loss

IV. BEST MODEL ARCHITECTURES

After evaluating the 60 models, the top **3 models** were selected based on their validation accuracy and were retrained on the entire training set (50,000 images).

1) *Activation Functions:* ReLU for models: CNN6, CNN2. And Leaky ReLU for CNN7.

2) *Optimizers:* While experimenting, the following optimizers were included: Adam, RMSProp, AdaGrad, and SGD. However, while selecting the model with the ray tune, it only picked models with **SGD**.

3) *Hyper-parameters for each model::*

- **CNN6:**

- Optimizer: SGD
- Learning Rate: 0.004278
- Momentum: 0.00958
- Dampening: 0.0946

- **CNN7:**

- Optimizer: SGD
- Learning Rate: 0.0390
- Momentum: 0.178957

- Dampening: 0.34405

- **CNN2:**

- Optimizer: SGD
- Learning Rate: 0.0498
- Momentum: 0.0867
- Dampening: 0.0415

Figure 2 shows the three best model architectures. It shows the different convolution layers used with their perspective outputs, the activation functions, different normalization, and generalization layers and the fully connected layers.

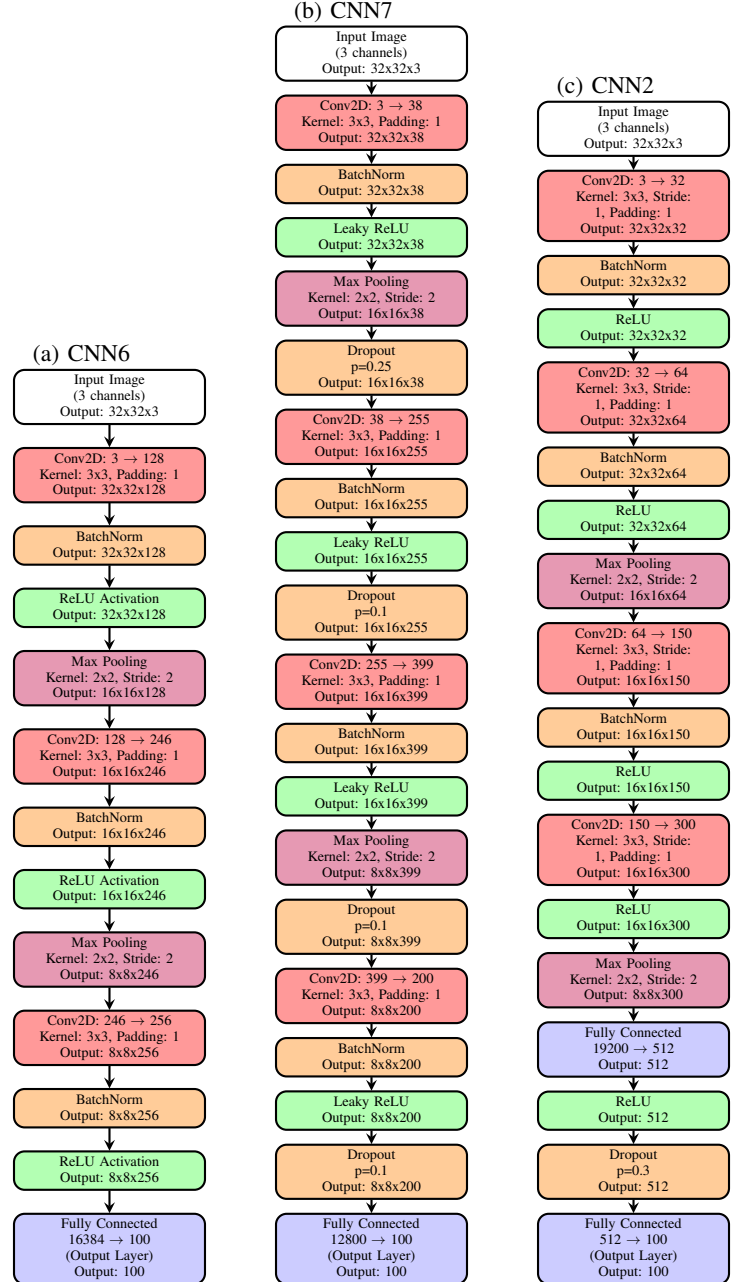


Fig. 2: Architectures of CNN6, CNN7, and CNN2

4) Model Selection Graphs:

V. FULL RETRAIN

After selecting the models, they were fully trained using the entire training set (50,000) images up to 100 epochs. Figure 3 shows the training loss of the 3 best models. Note that there is no validation loss because, at this point, all the training data was used for training.

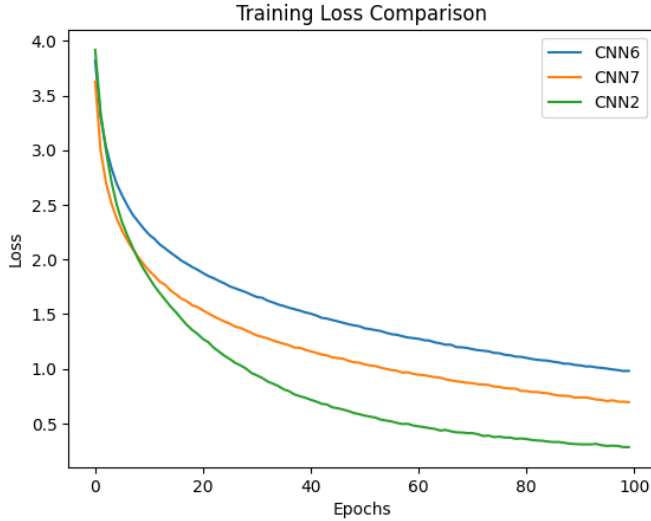


Fig. 3: Full Data Training

A. Test Accuracies

- CNN6: 60.1%
- CNN7: 61.1%
- CNN2: 58.7%

VI. COMPARISON WITH TOP MODELS

To compare these 3 best models, I used projects from [papers with code](#) that focus on training a model only on CIFAR100. The top 4 models found are

- [EXACT: How to Train Your Accuracy](#) - with an accuracy of 82.68
- [Differentiable Spike: Rethinking Gradient-Descent for Training Spiking Neural Networks](#) - with an accuracy of 74.24
- [Beta-Rank: A Robust Convolutional Filter Pruning Method For Imbalanced Medical Image Analysis](#) - with an accuracy of 74.01
- [Deep Residual Networks with Exponential Linear Unit](#) with an accuracy of 73.5

Taking an average of all three would give an average accuracy of 76.11%, and comparing it with the models developed here. The top models from papers with code perform about 16% better, which is comparable since I only focus on training this model for 1-2 weeks. The author of these competitive projects used more time to experiment and develop new ideas to improve their performance. Models that performed with 60% accuracy are ranked in the 180 positions, and with 50% accuracy are ranked in the 192 positions.

VII. TOTAL NUMBER OF PARAMETERS

- CNN6 Total Parameters: 2,494,022
- CNN7 Total Parameters: 3,004,917
- CNN2 Total Parameters: 10,393,946

VIII. CONCLUSION

The three CNN models developed performed well when trained solely on CIFAR100 data aided by data transformations, generalization and normalization layers, and ray tune for model selection.